Install hadoop in linux (ubentu):

Install JDK on Ubuntu:

$ sudo apt update

$ sudo apt install openjdk-8-jdk -y

- Once the installation process is complete, verify the current Java version:

$ java -version; javac -version

```
pnap@phoenixnap:~$ java -version; javac -version
openjdk version "1.8.0_422"
OpenJDK Runtime Environment (build 1.8.0_422-8u422-b05-1~24.04-b05)
OpenJDK 64-Bit Server VM (build 25.422-b05, mixed mode)
javac 1.8.0_422
```

Install OpenSSH on Ubuntu

$ sudo apt install openssh-server openssh-client -y

- In the example below, the output confirms that the latest version is already installed.

```
pnap@phoenixnap:~$ sudo apt install openssh-server openssh-client -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
openssh-server is already the newest version (1:9.6p1-3ubuntu13.5).
openssh-client is already the newest version (1:9.6p1-3ubuntu13.5).
openssh-client set to manually installed.
0 upgraded, 0 newly installed, 0 to remove and 166 not upgraded.
```

Create Hadoop User

$ sudo adduser hdoop

- The username, in this example, is **hdoop**. You are free to use any username and password you see fit.

- Switch to the newly created user and enter the corresponding password:

$ su – hdoop

Enable Passwordless SSH for Hadoop User

$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa

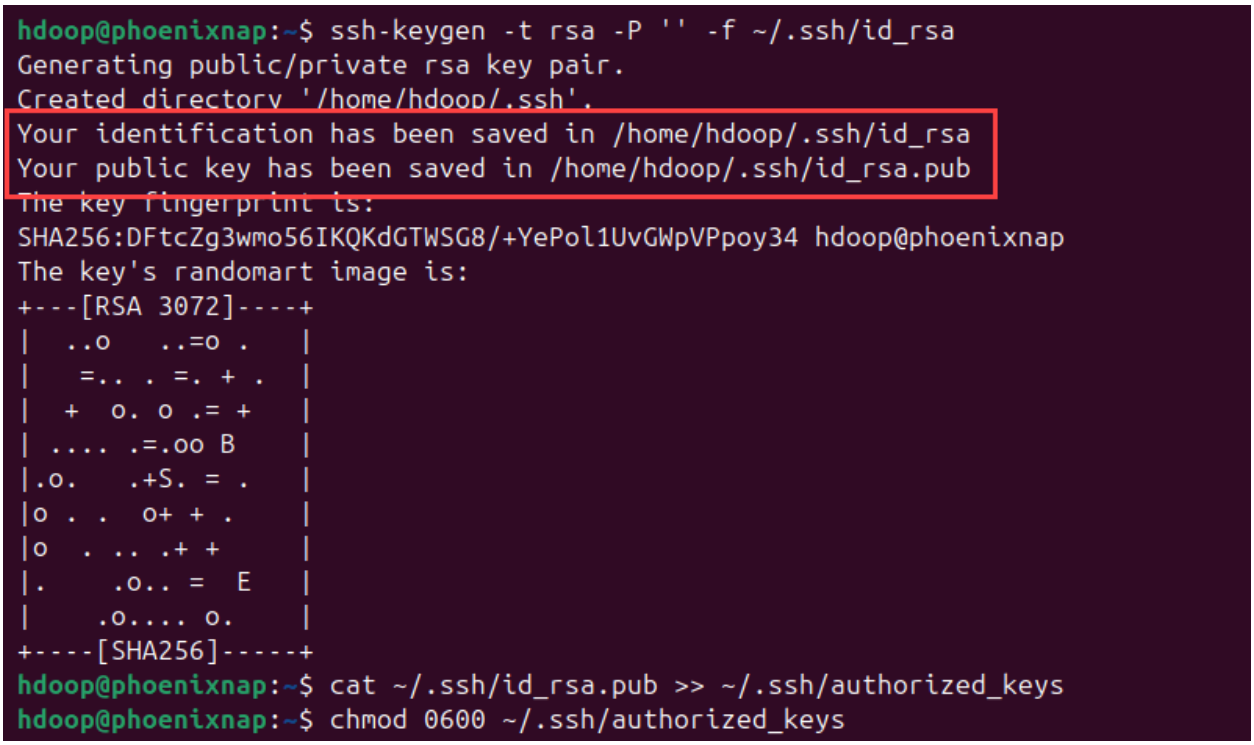

$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys

$ chmod 0600 ~/.ssh/authorized_keys

$ ssh localhost

Download and Install Hadoop on Ubuntu

$ wget https://dlcdn.apache.org/hadoop/common/hadoop-3.4.0/hadoop-3.4.0.tar.gz

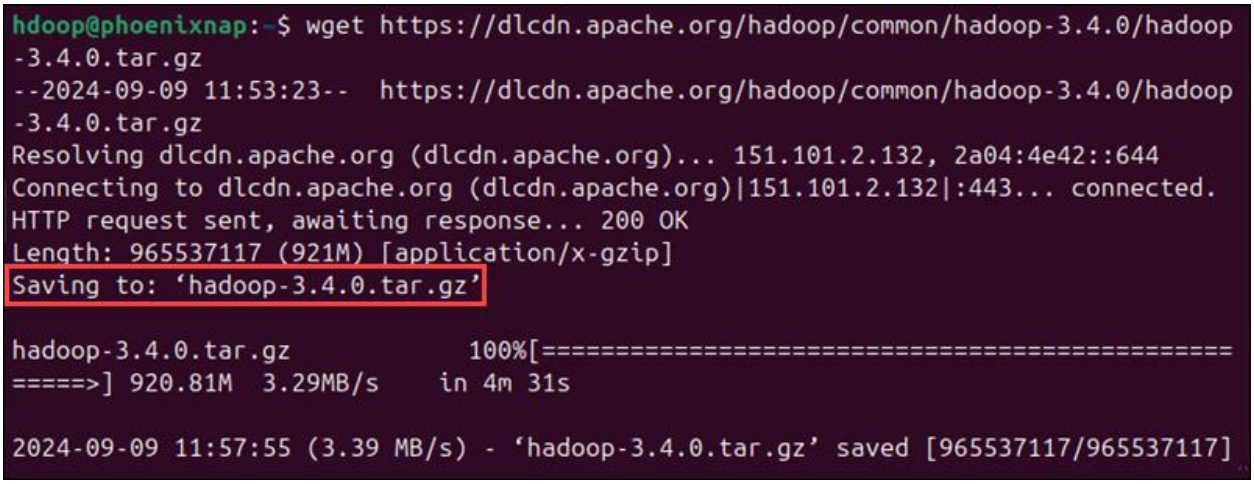

$ tar xzf hadoop-3.4.0.tar.gz

$nano .bashrc

-insert this :

```
#Hadoop Related Options
export HADOOP_HOME=/home/hdoop/hadoop-3.4.0
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native"
```

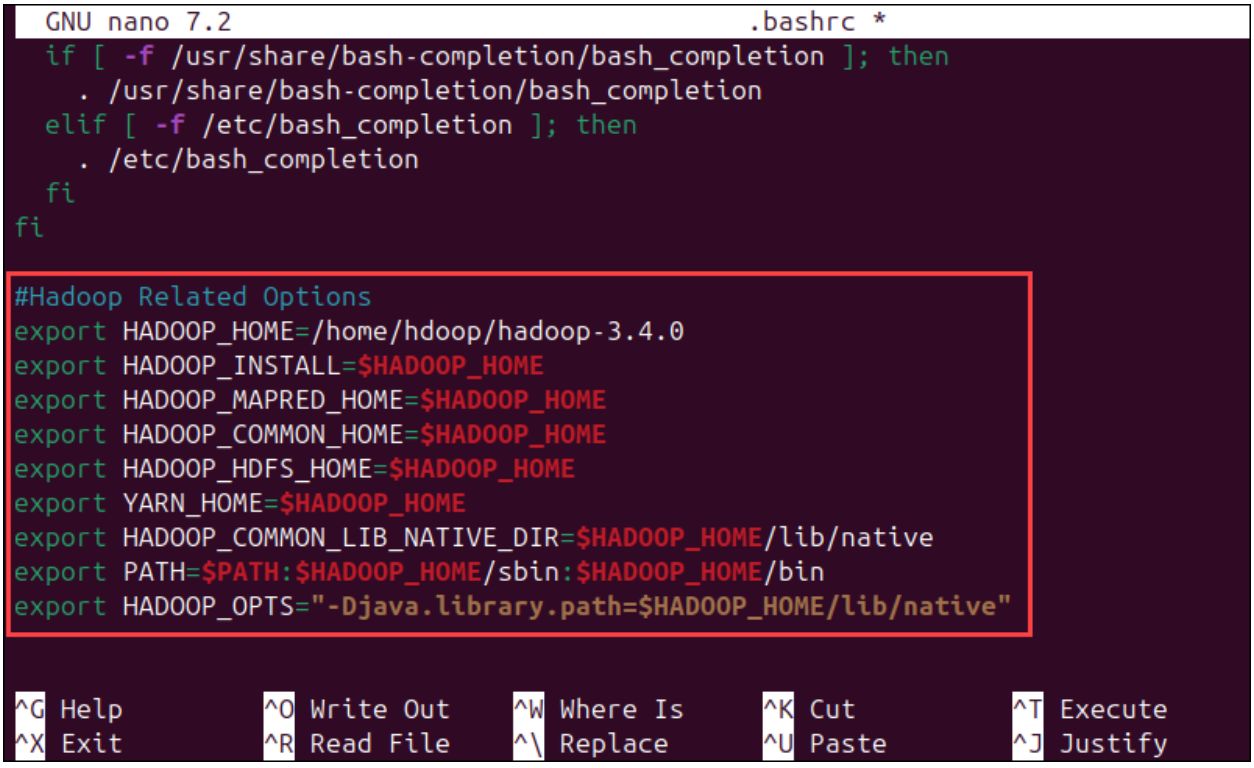

-click Ctrl+s and Ctrl+x to go out the editor

$source ~/.bashrc

#**Edit hadoop-env.sh File**

$nano $HADOOP_HOME/etc/hadoop/hadoop-env.sh

- Uncomment the **$JAVA_HOME** variable (i.e., remove the **#** sign) and add the full path to the OpenJDK installation on your system. If you have installed the same version as presented in the first part of this tutorial, add the following line:

$export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64

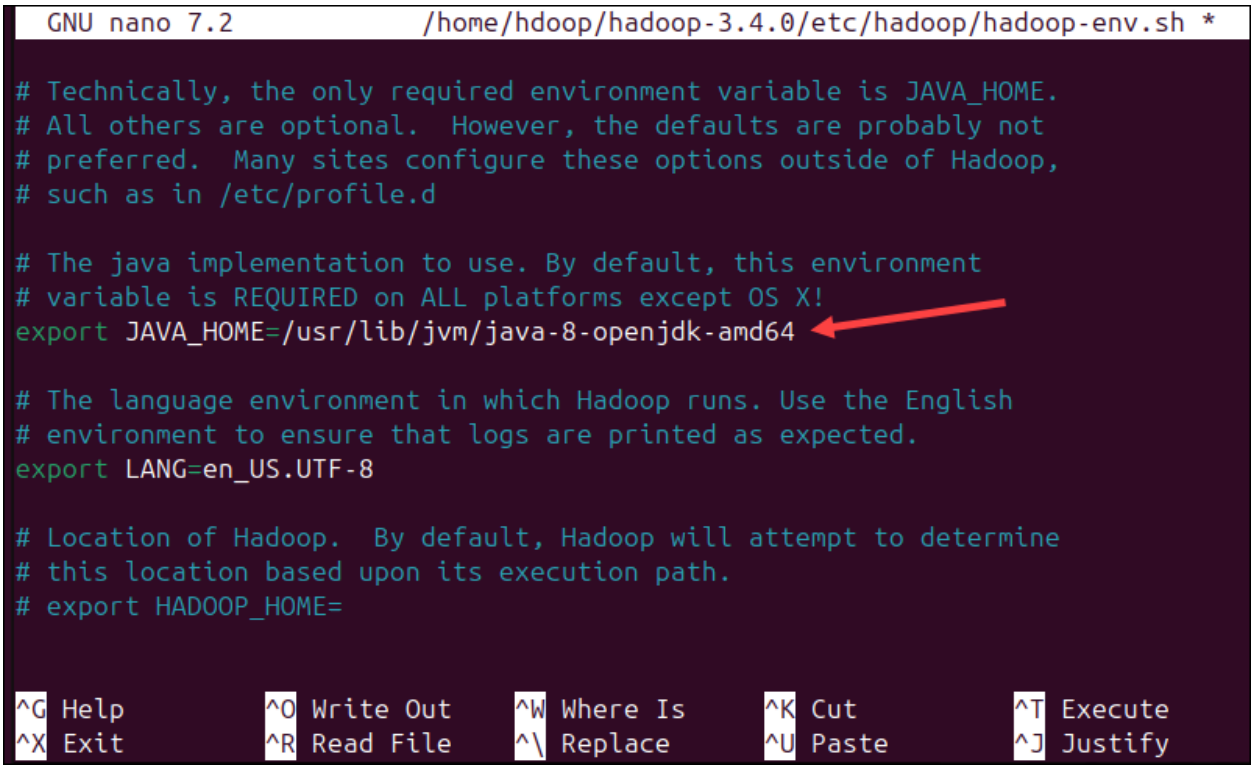

-click Ctrl+s and Ctrl+x to go out the editor

#Edit core-site.xml File

Note: keep this number , we will use it alot

$hostname -I | awk '{print $1}'

-copy the number address

$nano $HADOOP_HOME/etc/hadoop/core-site.xml

-Add the following configuration to override the default values for the temporary directory and add your HDFS URL to replace the default local file system setting:

```
<configuration>
<property>
 <name>hadoop.tmp.dir</name>
 <value>/home/hdoop/tmpdata</value>
```

</property>

<property>

<name>fs.default.name</name>

<value>hdfs://put the number here:9000</value>

</property>

</configuration>

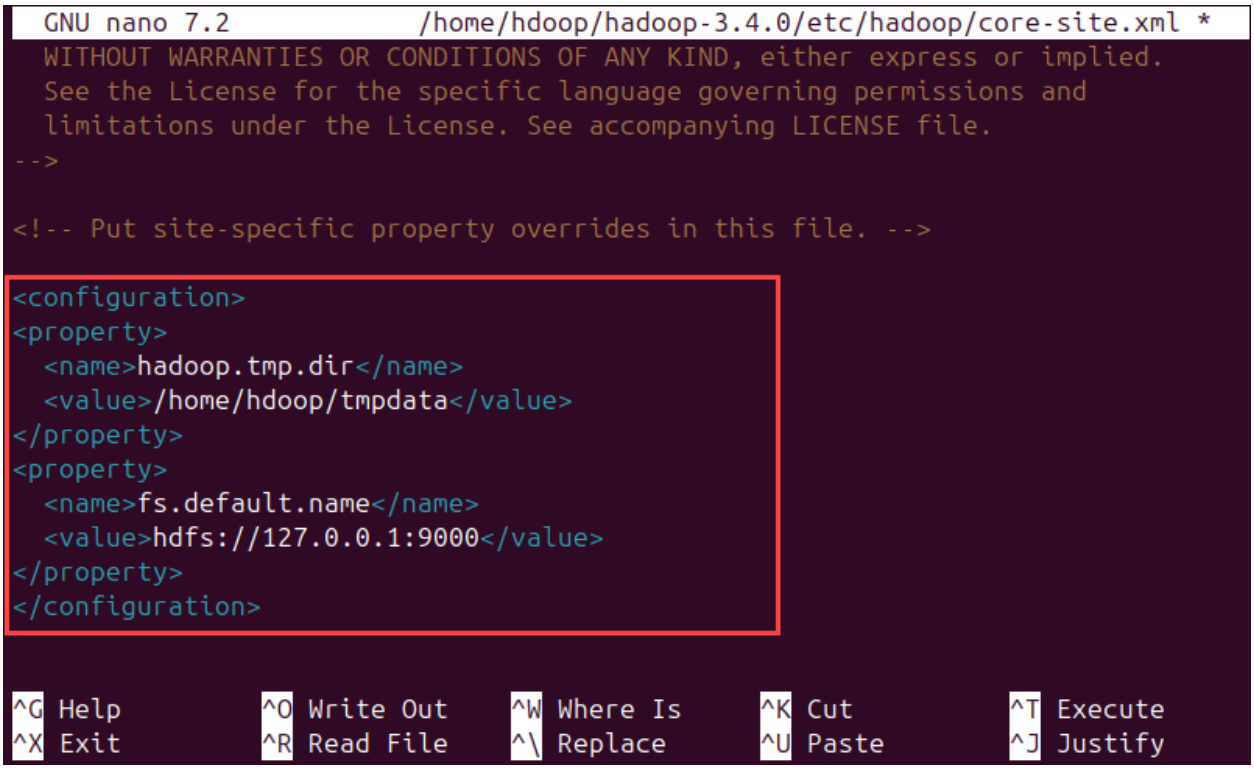

#Edit hdfs-site.xml File

$ sudo nano $HADOOP_HOME/etc/hadoop/hdfs-site.xml

-Add the following configuration to the file and, if needed, adjust the NameNode and DataNode directories to your custom locations:

<configuration>

<property>

```
 <name>dfs.data.dir</name>
 <value>/home/hdoop/dfsdata/namenode</value>
</property>
<property>
 <name>dfs.data.dir</name>
 <value>/home/hdoop/dfsdata/datanode</value>
</property>
<property>
 <name>dfs.replication</name>
 <value>1</value>
</property>
</configuration>
```

Note: -make sure you replace “hdoop” with your user, you can find it in the command line,like here :

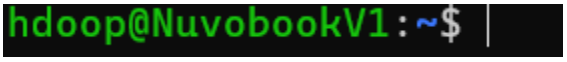

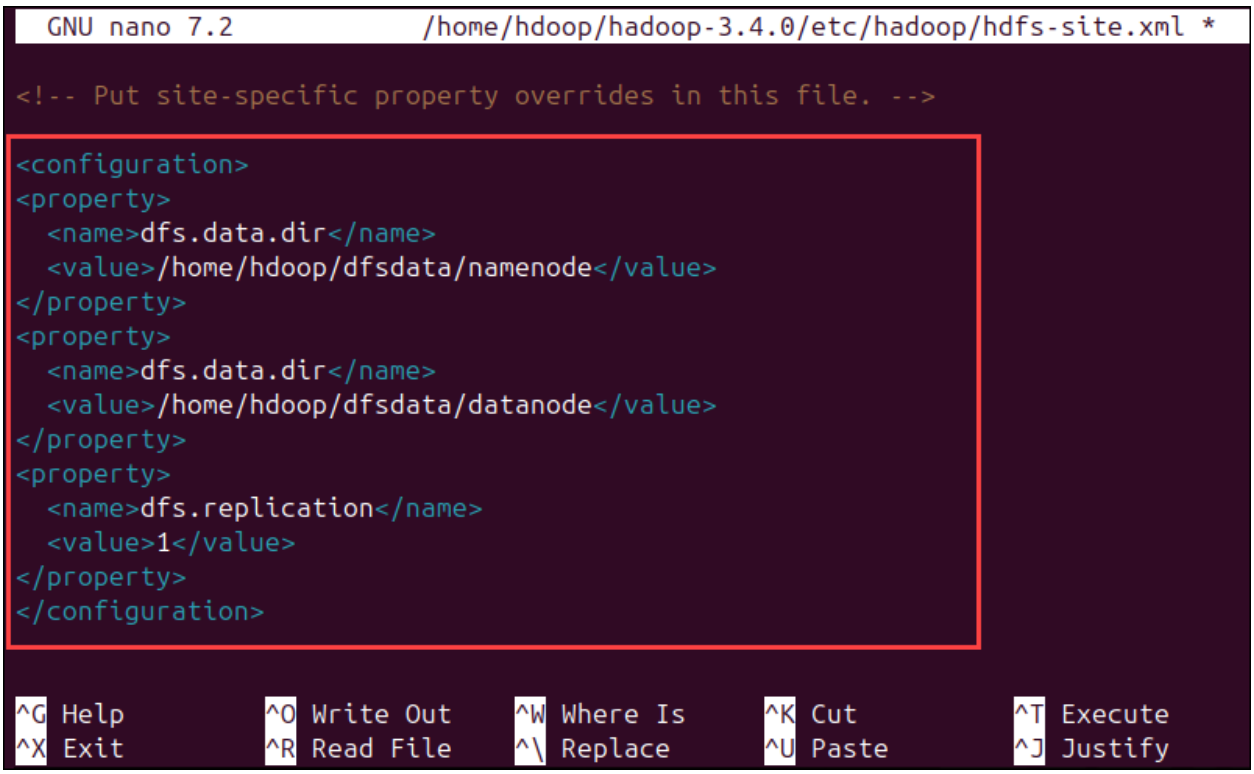

#Edit mapred-site.xml File

$sudo nano $HADOOP_HOME/etc/hadoop/mapred-site.xml

-Add the following configuration to change the default MapReduce framework name value to **yarn**:

```
<configuration>
<property>
 <name>mapreduce.framework.name</name>
 <value>yarn</value>
</property>
</configuration>
```

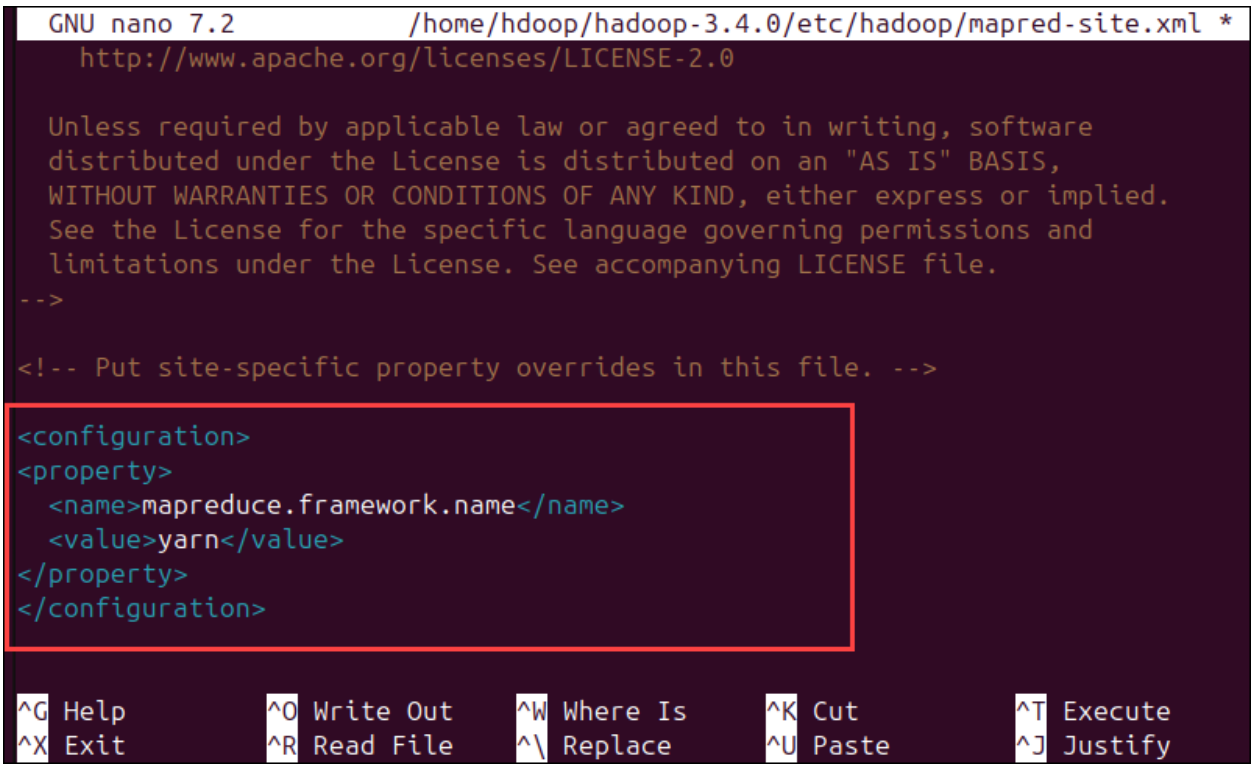

Edit yarn-site.xml File

$nano $HADOOP_HOME/etc/hadoop/yarn-site.xml

-Append the following configuration to the file:

```
<configuration>
<property>
 <name>yarn.nodemanager.aux-services</name>
 <value>mapreduce_shuffle</value>
</property>
<property>
 <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
 <value>org.apache.hadoop.mapred.ShuffleHandler</value>
```

</property>

<property>

<name>yarn.resourcemanager.hostname</name>

<value>put the same number that you puut before here</value>

</property>

<property>

<name>yarn.acl.enable</name>

<value>0</value>

</property>

<property>

<name>yarn.nodemanager.env-whitelist</name>

<value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PERPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>

</property>

</configuration>

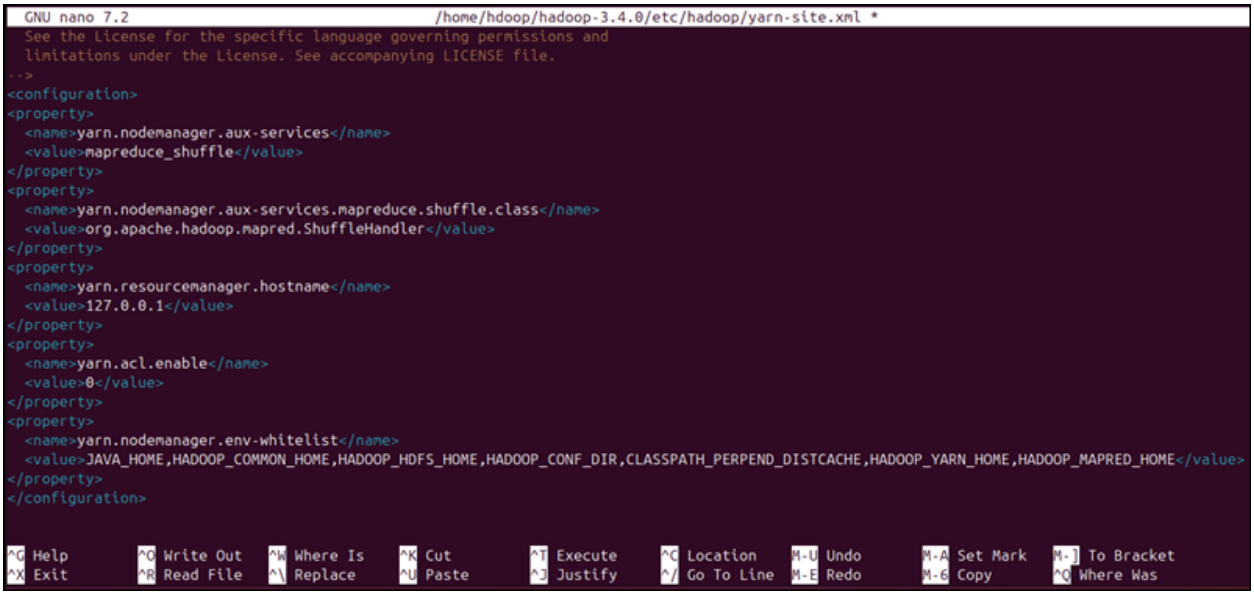

Format HDFS NameNode

$hdfs namenode -format

-The shutdown notification signifies the end of the NameNode format process.

```
hdoop@phoenixnap:~$ hdfs namenode -format
WARNING: /home/hdoop/hadoop-3.4.0/logs does not exist. Creating.
2024-09-09 13:08:42,739 INFO namenode.NameNode: STARTUP_MSG:
/************************************************************
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = phoenixnap/127.0.1.1
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 3.4.0
STARTUP_MSG:   classpath = /home/hdoop/hadoop-3.4.0/etc/hadoop:/home/hdoop/hadoo
3.4.0/share/hadoop/common/lib/curator-client-5.2.0.jar:/home/hdoop/hadoop-3.4.0/
2024-09-09 13:08:45,012 INFO namenode.FSNamesystem: Stopping services started fo
standby state
2024-09-09 13:08:45,018 INFO namenode.FSImage: FSImageSaver clean checkpoint: tx
=0 when meet shutdown.
2024-09-09 13:08:45,019 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at phoenixnap/127.0.1.1
************************************************************/
```

Start Hadoop Cluster

$cd

$cd *hadoop-3.4.0/sbin*

- execute the following command to start the NameNode and DataNode:

$./start-dfs.sh

```
hdoop@phoenixnap:~/hadoop-3.4.0/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [phoenixnap]
```

- Once the *namenode*, *datanodes*, and *secondary namenode* are up and running, start the YARN resource and *nodemanagers* by typing:

$./start-yarn.sh

```
hdoop@phoenixnap:~/hadoop-3.4.0/sbin$ ./start-yarn.sh
Starting resourcemanager
Starting nodemanagers
```

-Run the following command to check if all the daemons are active and running as Java processes:

$jps

-If everything works as intended, the resulting list of running Java processes contains all the HDFS and YARN daemons.

```
hdoop@phoenixnap:~/hadoop-3.4.0/sbin$ jps
45169 DataNode
46355 ResourceManager
45033 NameNode
46476 NodeManager
45373 SecondaryNameNode
47390 Jps
```

Access Hadoop from Browser

#nameNode

http://localhost:9870

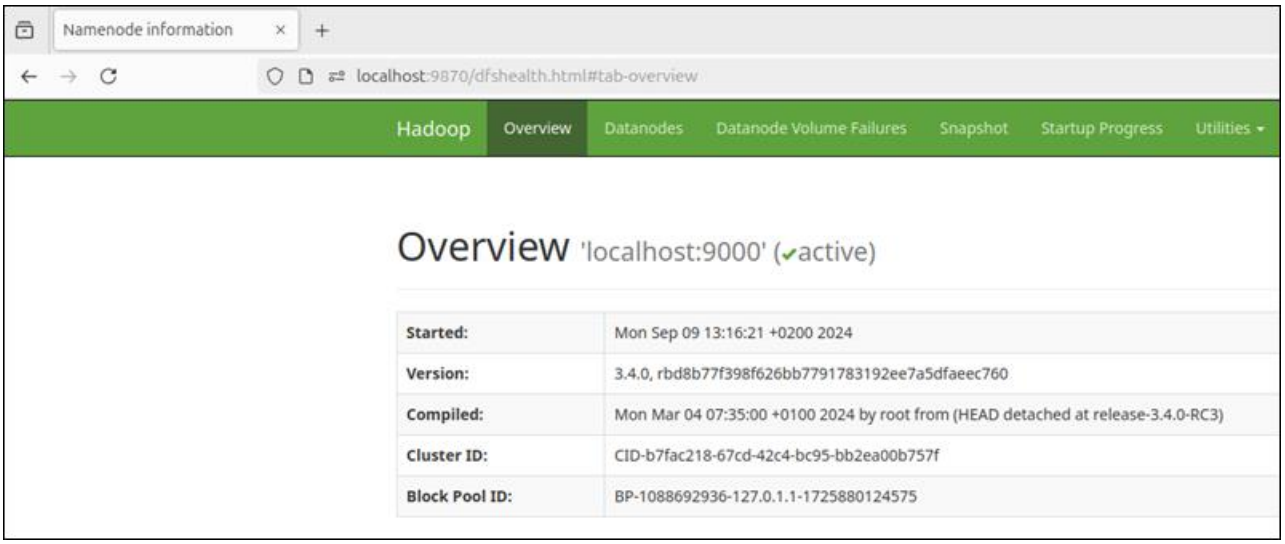

#dataNode:

http://localhost:9864

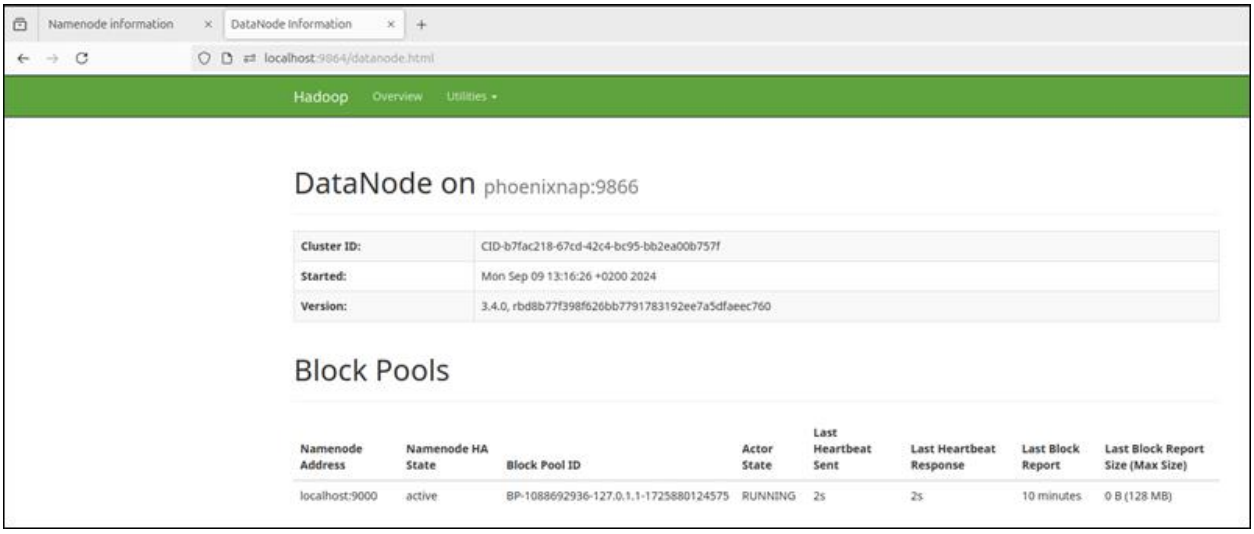

#resorce manager:

http:// put the same number here:8088